Advanced Hbase

Student Name: Subbham Vishal

Course: Big Data Hadoop & Spark Training

Assignment −
Execute blog present in below link

https://acadgild.com/blog/importtsv-data-from-hdfs-into-hbase/

Solution : The screenshots of the process is as follows

```
[acadgild@localhost ~]$ hbase shell
2016-11-02 15:59:31,727 INFO  [main] Configuration.deprecation: hadoop.native.li
b is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 0.98.14-hadoop2, r4e4aabb93b52f1b0fef6b66edd06ec8923014dec, Tue Aug 25 2
2:35:44 PDT 2015

hbase(main):001:0>
```

```
hbase(main):001:0> create 'bulktable','cf1','cf2'
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/
org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/li
b/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2016-11-02 16:01:01,494 WARN  [main] util.NativeCodeLoader: Unable to load nativ
e-hadoop library for your platform... using builtin-java classes where applicabl
e
0 row(s) in 2.7600 seconds

=> Hbase::Table - bulktable
hbase(main):002:0>
```

```
[acadgild@localhost ~]$ mkdir Hbase
[acadgild@localhost ~]$
[acadgild@localhost ~]$ cd Hbase
[acadgild@localhost Hbase]$
```

```
[acadgild@localhost Hbase]$ vi bulk_data.tsv
```

acadgild@localhost:~

File   Edit   View   Search   Terminal   Help

```
1       Amit    4
2       girja   3
3       jatin   5
4       swati   3
~
~
```

```
[acadgild@localhost Hbase]$ hadoop fs -mkdir /hbase
16/11/02 21:10:24 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
[acadgild@localhost Hbase]$ hadoop fs -put bulk_data.tsv /hbase/
16/11/02 21:10:48 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
[acadgild@localhost Hbase]$
16/11/02 21:11:34 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
1       Amit    4
2       girja   3
3       jatin   5
4       swati   3
[acadgild@localhost Hbase]$
```

```
[acadgild@localhost Hbase]$ hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.columns=HBASE_ROW_KEY,cf1:name,cf2:exp bulktable /hbase/bulk
_data.tsv
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop-2.6.0/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
2016-11-02 16:18:44,380 WARN  [main] util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where
 applicable
2016-11-02 16:18:44,494 INFO  [main] zookeeper.RecoverableZooKeeper: Process identifier=hconnection-0x6572421 connecting to ZooKeeper ensemble=localho
st:2181
2016-11-02 16:18:44,507 INFO  [main] zookeeper.ZooKeeper: Client environment:zookeeper.version=3.4.6-1569965, built on 02/20/2014 09:09 GMT
2016-11-02 16:18:44,507 INFO  [main] zookeeper.ZooKeeper: Client environment:host.name=localhost
2016-11-02 16:18:44,507 INFO  [main] zookeeper.ZooKeeper: Client environment:java.version=1.8.0_65
2016-11-02 16:18:44,507 INFO  [main] zookeeper.ZooKeeper: Client environment:java.vendor=Oracle Corporation
2016-11-02 16:18:44,507 INFO  [main] zookeeper.ZooKeeper: Client environment:java.home=/usr/local/java/jre
2016-11-02 16:18:44,520 INFO  [main] zookeeper.ZooKeeper: Client environment:java.class.path=/usr/local/hbase/conf:/usr/local/java/lib/tools.jar:/usr/
local/hbase:/usr/local/hbase/lib/activation-1.1.jar:/usr/local/hbase/lib/aopalliance-1.0.jar:/usr/local/hbase/lib/asm-3.1.jar:/usr/local/hbase/lib/avr
o-1.7.4.jar:/usr/local/hbase/lib/commons-beanutils-1.7.0.jar:/usr/local/hbase/lib/commons-beanutils-core-1.8.0.jar:/usr/local/hbase/lib/commons-cli-1.
2.jar:/usr/local/hbase/lib/commons-codec-1.7.jar:/usr/local/hbase/lib/commons-collections-3.2.1.jar:/usr/local/hbase/lib/commons-compress-1.4.1.jar:/u
sr/local/hbase/lib/commons-configuration-1.6.jar:/usr/local/hbase/lib/commons-daemon-1.0.13.jar:/usr/local/hbase/lib/commons-digester-1.8.jar:/usr/loc
```

```
2016-11-02 16:18:50,713 INFO  [main] Configuration.deprecation: mapred.cache.files.filesizes is deprecated. Instead, use mapreduce.job.cache.files.fil
esizes
2016-11-02 16:18:50,713 INFO  [main] Configuration.deprecation: fs.default.name is deprecated. Instead, use fs.defaultFS
2016-11-02 16:18:51,197 INFO  [main] mapreduce.JobSubmitter: Submitting tokens for job: job_1478096881917_0003
2016-11-02 16:18:52,372 INFO  [main] impl.YarnClientImpl: Submitted application application_1478096881917_0003 to ResourceManager at /0.0.0.0:8032
2016-11-02 16:18:52,525 INFO  [main] mapreduce.Job: The url to track the job: http://http://localhost:8088/proxy/application_1478096881917_0003/
2016-11-02 16:18:52,526 INFO  [main] mapreduce.Job: Running job: job_1478096881917_0003
2016-11-02 16:19:41,977 INFO  [main] mapreduce.Job: Job job_1478096881917_0003 running in uber mode : false
2016-11-02 16:19:42,131 INFO  [main] mapreduce.Job:  map 0% reduce 0%
2016-11-02 16:20:50,804 INFO  [main] mapreduce.Job:  map 100% reduce 0%
2016-11-02 16:21:17,653 INFO  [main] mapreduce.Job: Job job_1478096881917_0003 completed successfully
Exception in thread "main" java.lang.IllegalArgumentException: No enum constant org.apache.hadoop.mapreduce.JobCounter.MB_MILLIS_MAPS
        at java.lang.Enum.valueOf(Enum.java:238)
        at org.apache.hadoop.mapreduce.counters.FrameworkCounterGroup.valueOf(FrameworkCounterGroup.java:148)
        at org.apache.hadoop.mapreduce.counters.FrameworkCounterGroup.findCounter(FrameworkCounterGroup.java:182)
        at org.apache.hadoop.mapreduce.counters.AbstractCounters.findCounter(AbstractCounters.java:154)
        at org.apache.hadoop.mapreduce.TypeConverter.fromYarn(TypeConverter.java:240)
        at org.apache.hadoop.mapred.ClientServiceDelegate.getJobCounters(ClientServiceDelegate.java:370)
        at org.apache.hadoop.mapred.YARNRunner.getJobCounters(YARNRunner.java:511)
        at org.apache.hadoop.mapreduce.Job$7.run(Job.java:756)
        at org.apache.hadoop.mapreduce.Job$7.run(Job.java:753)
        at java.security.AccessController.doPrivileged(Native Method)
        at javax.security.auth.Subject.doAs(Subject.java:422)
        at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1491)
        at org.apache.hadoop.mapreduce.Job.getCounters(Job.java:753)
        at org.apache.hadoop.mapreduce.Job.monitorAndPrintJob(Job.java:1361)
        at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1289)
        at org.apache.hadoop.hbase.mapreduce.ImportTsv.run(ImportTsv.java:680)
        at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:70)
        at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:84)
        at org.apache.hadoop.hbase.mapreduce.ImportTsv.main(ImportTsv.java:684)
[acadgild@localhost Hbase]$
```

```
hbase(main):003:0> scan 'bulktable'
ROW                      COLUMN+CELL
 1                       column=cf1:name, timestamp=1478083723895, value=Amit
 1                       column=cf2:exp, timestamp=1478083723895, value=4
 2                       column=cf1:name, timestamp=1478083723895, value=girja
 2                       column=cf2:exp, timestamp=1478083723895, value=3
 3                       column=cf1:name, timestamp=1478083723895, value=jatin
 3                       column=cf2:exp, timestamp=1478083723895, value=5
 4                       column=cf1:name, timestamp=1478083723895, value=swati
 4                       column=cf2:exp, timestamp=1478083723895, value=3
4 row(s) in 1.4550 seconds

hbase(main):004:0>
```